



Script for presenting the weekly update

For me only. The document on the table is PDQ_WEEKLY_UPDATE_26_SEPTMBER_2026.pdf.

Don't read this word for word or it will sound rehearsed. Read a section, look up, and say it the way it comes out. The numbers are the only part to get exactly right.

About seven minutes at a normal pace.

Opening

Have the document closed. Just talk.

Before I open the document, here's what my week actually looked like, because it was really four jobs.

First I went and got the clips. I found out how the pod downloads TikTok, re-pulled the lists, and went from 39 clips to 125, which finally gave me real evidence to test with.

Second, I took last week's rule change and measured it properly against the pod's original rules, on every clip I have, instead of the one clip I had before.

Third, I made two more rule changes, because with more clips I could see two other rules making the same mistake, and I measured those the same way.

And fourth, I worked out which of the remaining wrong answers can be fixed with rules and which ones can't, so the ones that can't go to the right place instead of me chasing them forever.

The headline is: two wrong answers stopped, zero correct answers broken, and I can prove both of those numbers. I'll walk through it in order.

1. Getting the clips back

Open the document at section 1. Point at the small table with 80, 141, 39, 125.

The first thing I did was go and get clips, because you told me last week to test with clips where I know the right answer, and I only had a handful.

The way I did it was mostly by reading our own code. Instead of asking again how the pod downloads TikTok, I opened the pod's downloader and found it falls back to the ScrapeCreators API when yt-dlp gets blocked, using the token that's already in the credentials file. So I called the same function the pod calls. I also re-pulled the two lists from the database, and I rebuilt my index (the library the engine searches) so every film on those lists is in it.

That took me from 39 downloaded clips to 125, and from about 12 I could actually test with to 58.

And there's one finding in here I want to flag on its own. The clips rot. Fourteen are deleted from the platforms permanently, and of the older clips only three in ten still match what was traced. The fresh ones survive at six in ten. So every week we wait, Trace loses a bit more of its ability to learn from its own mistakes. I'll come back to that at the end because I think there's a cheap fix.

2. Measuring my change properly

Point at section 2, then hold a finger on the words "nothing else moves".

Next I went back and properly tested the rule change I made last week, the P8 and P6 one that stopped the Demeter clip.

The way I test now is that I keep the pod's original rules and my rules side by side, and a script runs every single clip through both and compares. The clips users confirmed as correct are the safety net. If my version loses any of them, the change is no good, however many wrong answers it removes.

The whole job is doing it without breaking answers that already work, and last week I could only check that on one clip, so I believed my change was safe but I couldn't prove it.

Now I can. Every correct answer that works on the pod's rules also works on mine.

For Trace that means the fix is real. One wrong film name that was reaching someone's phone doesn't any more, and nobody who was getting the right answer loses it.

3. What a real match gives

Point at section 3, at the numbers 24.8 and 38.0, then at 41.4.

Quick reminder of how scoring works, because this section is about where to draw a line. Every match gets a score for how different the two pictures are. Zero means identical, so lower is better, and anything worse than 45 gets thrown away as meaningless.

My change last week drew a line at 40. Evidence scoring worse than 40 can't decide an answer any more. That's what blocked Demeter, which scored 41.4.

What I didn't know when I put that limit is where the correct answers score. So this week I measured it. I took the clips users confirmed as right and wrote down what the correct film scored on each one. They come in between about 25 and 38.

Put those together and my line at 40 sits above every correct answer and below the wrong one, which is the only place a line works. But that is still narrow. The worst correct answer is 38 and the wrong one is 41.4, so the safe zone is about three points wide, and I picked 40 before I knew any of this. If I'd picked 37 instead, I would have blocked two correct answers and never understood why.

So the lesson I'm taking is measure first, draw the line second. From here on the thresholds in this engine come from what real matches actually score, and that's what protects the answers users are already happy with.

4. The idea I tried and dropped

Point at section 4. This one I present as a win, not an apology.

I also want to tell you about something that didn't work, because it's the best example of the method doing its job.

While testing the correct clips I found the engine gives up on a clip after about a second and a half if nothing has matched yet, no matter how long the clip is. One clip had the film in it at distance 12, which is basically the same picture, and the engine quit at 1.88 seconds of a 57-second video. The pre-scan is the only reason clips like that get found.

So I tried the obvious fix, which was to make it look at a quarter of the clip before giving up. When I first tried it on just four clips, it looked great. Then I ran the full measurement on both lists, and it let two wrong answers back in for every correct one it recovered. Looking longer helps real matches, but it also gives look-alikes time to pile up frames until a loose rule fires.

So I dropped it. For Trace, the impact is a bad trade that never shipped. A version of the app that recovers one right answer but sends out two extra wrong ones is worse, and the only reason I know that is because I measured before keeping it.

5. Most answers never reach the rules

Point at the two-row table in section 5.

Then I asked a question I couldn't ask before I had enough clips, which is: which part of the engine actually produces the answers?

I counted it. The pre-scan, the quick skim that runs before any rule, decides about three quarters of everything, the right answers and the wrong ones alike. So all the work I've been doing, changing the numbers inside the rules, only ever affects the quarter of clips the pre-scan didn't already decide.

That matters for where we spend effort. Four of the remaining five wrong answers come from the pre-scan, so changing rule numbers can never touch them. But I'm not proposing we tighten the pre-scan either, because on my count it produces fifteen correct answers for every three or four wrong ones. It's the most accurate part of the engine. Squeezing it would put more good answers at risk than bad ones.

The useful thing is knowing the limit. We now know the most that changing these numbers can ever fix, so we can stop there instead of spending more weeks on changes that can't reach the remaining five.

6. Two more changes, and the final measurement

****Point at the three-row table in section 6. Finger down the right-hand column: 49, 49,**

49.**

With more clips I noticed two other rules doing the same thing P8 did with Demeter. They were deciding answers on far too little evidence. One of them answered Grave Encounters 2 with just two frames, and neither of them was even a close match. The other one, the probe, was allowed to answer with one single

frame.

To fix them I didn't invent anything new. I copied what the engine already does in other places. There's already a rule that demands more frames when the frames are lower quality, so I made the Grave Encounters rule work the same way. Two frames is enough only when they're good matches. If they're weaker, it needs four. And there's already a rule that won't accept one near-perfect frame on its own, it wants a second frame to back it up. So the probe works like that now too. It needs two frames pointing at the same film before it's allowed to decide.

Then I ran the big test. I took all forty wrong clips and all eighty-five confirmed correct clips, and ran them through three versions of the engine. The pod's version, my version from last week, and my version from this week.

The pod's version gives seven wrong answers. Last week's gives six. This week's gives five. And the correct answers stayed at forty-nine in all three versions, so I removed two wrong answers and didn't lose a single thing.

And I want to be clear about what stopping a wrong answer actually looks like. On those two clips, Demeter and Grave Encounters 2, the engine doesn't name a different film now. It stays quiet. The film still gets passed along as a hint, so audio and SSCD can still look at it, but PDQ no longer sends the wrong name on its own. That's the trade you told me to make, a no answer instead of a wrong one.

So for Trace, two wrong film names that were reaching users have stopped, and I checked that it cost us nothing.

7. Where this ends

Point at section 7. Then look up, because this is the conclusion.

The five wrong answers that are left are the reason I can say where this work ends.

I looked at each one. They match at distances between 6 and 22, and my correct answers match between 8 and 26. So the wrong ones score as well as the right ones, sometimes better. On the evidence PDQ sees, these five look exactly like correct answers. That's the proven part: any rule strict enough to block them would block real matches too.

What I can't tell you yet is why users pressed wrong on them, because for these five nobody typed in what the film really was. It could be clips that mix footage from more than one thing. It could be shots that turn up in more than one film, like trailer footage. And in at least one case the vote itself is shaky, because The In Between was voted correct by two users and wrong by one.

The obvious next step is to watch these five clips next to the scene the engine matched, and I have all five downloaded, so that's quick to do. But whichever way that comes out, it isn't a thresholds problem, and that's why I want to hand these over with their scorecards rather than keep chasing them with rule changes.

8. What I'd like from the meeting

Point at the three numbered asks. Say each one, then pause and let him respond.

So I need three things.

First, a yes or no on the two new rule changes. They're small, they're commented, and they're measured at zero cost. With last week's changes the whole diff is four thresholds and two logic changes.

Second, a home for those five picture-match cases, because they're a policy question now, not a tuning one.

And third, the one I'd push hardest. There's a table where the pod already stores every clip's fingerprints, and it deletes them after about a day. If we kept those rows for any clip a user votes on, we could replay any wrong answer forever without needing the video. That would end the clip rot problem permanently, and it's a retention change to something that already exists, not new infrastructure.

Close

That's the week. Two wrong answers gone, nothing broken, and a clear line between what's an engine problem and what isn't.

If a question throws me

The fallback line is the same as always: two wrong answers removed, zero right answers lost, and the other five aren't a rules problem. The longer answers are in MEETING_PREP.md and the numbers are all in the document itself.