



## MEETING SCRIPT

# Script for presenting the 4 October update

For me only. The document on the table is PDQ\_WEEKLY\_UPDATE\_04\_OCTOBER\_2026.pdf.

Same rule as last time: don't read it word for word. Read a section, look up, say it the way it comes out. The numbers are the only part to get exactly right.

About eight minutes at a normal pace.

## Opening

**Have the document closed. Just talk.**

Last week you told me to stop bending the thresholds and change the actual tech. So this week was three jobs.

First, I made the testing honest: every clip I run now records which rule gave the answer. So when something is wrong, I know exactly which part of the engine did it, not just that it happened.

Second, I compared the three parts of the engine that can answer, worked out what the careful rules have that the pre-scan and the probe don't, and rebuilt the probe out of the missing pieces.

Third, I built an edits test. Ten things a reposter actually does to a clip, applied one at a time to clips the engine gets right, so any failure is the edit's fault and nothing else. That found two places where the engine is completely blind, and I fixed both.

The headline: the pod's rules give seven wrong answers on my full set. The rebuilt probe with the rules behind it gives one. And of the fifty-two correct answers, I keep forty-eight. I'll go through it.

## 1. Which rule gave the wrong answers

**Open at section 1. Point at the screenshot.**

This is the printout, not a summary. All 125 clips through the pod's exact rules, and every answer tagged with the rule that produced it.

Seven wrong answers reproduce. Four of them come from the pre-scan, and look at the distances: six, six, twelve, eighteen. Those pictures genuinely are in the index. The rules never even ran on those clips, which is why the last two weeks of rule changes could never have caught them. The other three are P0, P9E and P8, which are exactly the three my earlier changes went after.

So the thing I'd been tuning could only ever reach three of the seven. That's the clearest argument I have for what you said last week.

## 2. What the system has that the deciders don't

### Section 2. The small table.

The engine can answer from three places. The pre-scan, the probe, and the rules.

The careful rules check three things: lots of frames, the frames agreeing on where in the film they land, and the frames coming from more than one moment of the clip.

The pre-scan and the probe, which between them decide almost everything, check none of that. The probe answers off one frame. And the engine's own notes say a true match gives the same offset for every frame. That's the strongest signal in the data and the two busiest parts of the engine ignore it.

## 3. The rebuilt probe

### Section 3. Don't linger, the next section is the point.

So the new probe is just: the pre-scan's whole-clip view, plus all the checks it never had. It samples forty frames across the clip, keeps a scorecard per film, and it only answers on three bars, each copied from a rule the pod already trusts. Two near-perfect frames. Or the old pre-scan bar plus agreement on film time. Or five medium frames that agree and come from two moments of the clip.

The old one-frame probe is retired. This replaces it.

## 4. The measurement

### Section 4. Point at the three-line table, then the screenshot.

Same yardstick as every week, both numbers, all 125 clips.

Pod's rules: seven wrong, fifty-two correct. The new probe entirely on its own: one wrong, forty-one. With the rules running behind it as the fallback: still one wrong, and forty-eight correct.

The one that survives is Krrish 3, at distance six. The pictures are identical to pictures in the index. Nothing picture-based can refuse that; it needs audio or SSCD.

The four correct answers I lose are all montage clips. The right film was the top candidate every single time, it just scatters across film time because the clip cuts between scenes. I can get them back, but the exact same loosening lets The Northman and Civil War back in, because on picture evidence they look the same. You told me a no answer beats a wrong one, so I left them out. But that's a policy call and I'd rather you made it than me.

One more thing from the tables. The clip the pod answered as The Misfits comes back as F1: The Movie, at distance eight, on both the pod's rules and mine. Worth watching that clip. Something doesn't add up about where the live answer came from.

## 5. The edits test

### Section 5. The edit table screenshot, then the before-and-after one.

Ten edits, each applied to eight clips the engine identifies correctly unedited.

First thing: no edit ever produced a wrong film. Anywhere. Edits don't fool the engine, they blind it. It goes quiet instead of guessing.

Filters, black and white, recompression, speed changes: barely a scratch. Rotation and zoom hurt. And two edits blinded it completely. Mirroring: zero out of eight. And letterboxing: one out of eight.

Both turned out to be missing tech. Nothing in the engine ever hashes a flipped frame, so I made the probe hash the mirror of everything it samples. Zero out of eight became eight out of eight, matching at distance two. And the pod has an active-region crop that finds the picture inside black bars, which is one of the pieces the laptop port left out. I built a simple version. One out of eight became six out of eight.

Then the safety gate: I re-ran all 125 real clips with the new views on. Not one answer changed. The extra views only add, they never subtract.

The honest cost is the probe is about four times slower with the extra views. On my laptop, irrelevant. If this goes near the pod, you'd run plain views first and only pay for the second look when the plain views find nothing.

## 6. Housekeeping

### Section 6. Quick.

Three small things. Everything is in a private GitHub repo now, pdq-testing, code and results, and I can add you to it today.

Pod access works from my machine, verified read-only, nothing touched.

And while setting that up, the access doc with the key in it ended up pasted into a chat. So if you'd rather rotate the key, now's a clean moment. Your call.

## 7. The asks

### Close the document. Count on fingers.

Four things.

One, the policy call on the four montage clips. I recommend keeping the strict trade.

Two, someone watches the Misfits clip.

Three, a direction on the probe. If you want it taken towards the pod as the pre-scan and P0 replacement, I'd start with the plain-views-first version and measure it at the pod's index size.

And four, still open from last week: keeping the clip\_query\_hashes rows for voted clips. The clips are still rotting.

## If he asks...

**...why not loosen the probe to get the four montages back?** > Because I tried the numbers. The four I'd recover and the two wrong ones I'd readmit > sit in the same evidence range. There's no bar between them. It's a trade, not a fix.

**...whether the mirror views could cause new wrong answers?** > That was my worry too, so that was the gate. All 125 real clips, before and after, > zero answers changed. It's in the repo if you want the diff.

**...how long the probe takes now?** > About two minutes a clip on my laptop with every view on, against thirty seconds > before. Plain-views-first brings ordinary clips back to today's cost.

**...whether the speed edits broke the film-time agreement?** > I expected them to and they didn't. The three-second tolerance absorbs a 25 percent > speed change. I had a fix designed and threw it away because the measurement said it > wasn't needed.