



MEETING SCRIPT

Script for presenting the 4 October update

For me only. The document on the table is PDQ_WEEKLY_UPDATE_04_OCTOBER_2026.pdf.

Same rule as last time: don't read it word for word. Read a section, look up, say it the way it comes out. The numbers are the only part to get exactly right.

About nine minutes at a normal pace.

Opening

Have the document closed. Just talk.

Last week you told me to change the actual tech instead of just tightening the rules. So this week I did three things.

First, I changed how I test. Before, I could tell you a clip got the wrong film, but not which part of the engine made the mistake. Now, every time I run a clip, it writes down which rule gave the answer. So every mistake has a name next to it.

Second, I went looking for what the careful parts of the engine have that the quick parts don't, and I rebuilt the probe out of what was missing.

Third, I tested edited clips. I took clips the engine gets right, made ten edited copies of each, the kinds of edits people actually do when they repost, and ran those. If a copy fails, the edit is the reason, because the original worked. That found two edits where the engine is completely blind, and I fixed both.

The numbers first: on my full test set, the pod's rules give seven wrong answers. My new version gives one. And of the fifty-two clips the pod gets right, I still get forty-eight. I'll go through it.

1. Which rule gave the wrong answers

Open at section 1. Point at the screenshot.

This is the printout. All 125 clips through the pod's exact rules, and every answer tagged with the rule that produced it.

Seven wrong answers show up, and four of them come from the pre-scan. The pre-scan is the quick skim the engine does before anything else: it grabs forty frames spread across the clip, and if some film clearly has those pictures, it answers on the spot and nothing else ever runs.

Now look at the distances on those four. Distance is how different two pictures are. Zero means identical, and anything under about twenty is the same picture for all practical purposes. These four matched at six, six, twelve and eighteen. So those pictures genuinely are in our library, just attached to

the wrong film, probably shared footage, things like trailers. The rules never ran on those clips at all. Which means two weeks of me adjusting rules could never have fixed them, whatever numbers I picked.

The other three wrong answers did come from rules, and they're exactly the three rules I'd already changed. So the tuning I was doing could only ever reach three of the seven. That's the clearest proof of what you told me last week.

2. What the careful rules have that the quick parts don't

Section 2. The small table.

The engine can answer from three places: the pre-scan, the probe, and the rules. The strange thing is how differently careful they are.

The careful rules check three things before they answer. Did a lot of frames match. Do those frames all point at the same place in the film. And did they come from different moments of my clip, not just one frozen frame.

That middle check is worth thirty seconds, because the whole week hangs on it. If my clip really is a scene from a film, then every frame of my clip matches the film at the same point. First frame matches minute forty, next frame matches minute forty plus one second, and so on. They all agree. A coincidence doesn't look like that. A coincidence picks up one stray frame from minute ten, another from minute fifty, scattered all over. The engine's own documentation says this is what a true match looks like.

But the pre-scan and the probe, which between them give most of the answers, do none of these checks. The probe would answer off one single frame. So the parts doing most of the deciding do the least checking.

3. The rebuilt probe

Section 3. Don't linger, the next section is the point.

The new probe is those missing pieces put together. It skims the whole clip, forty frames, like the pre-scan. But instead of just counting hits, it keeps a scorecard for every film: how many frames hit it, how good the best one is, whether they agree on where in the film they land, and whether they came from different parts of my clip.

And it only answers in three situations, each copied from a rule the pod already trusts. Two near-identical frames. Or what the pre-scan used to need, plus the frames agreeing on film position. Or five decent frames that agree and come from two different moments of the clip.

The old probe, the one-frame one, is switched off. This replaces it.

4. The measurement

Section 4. Point at the three-line table, then the screenshot.

Both numbers, same as every week, all 125 clips.

The pod's rules: seven wrong, fifty-two right. The new probe completely on its own, everything else switched off: one wrong, forty-one right. And the version I'd actually propose, the new probe with the rules kept as a backup for the clips it isn't sure about: one wrong, forty-eight right.

The one wrong answer left is Krrish 3, matching at distance six. The pictures in that clip are in our library, and they are identical. No check that works on pictures can turn that away. That one needs the audio detective or SSCD.

The four right answers I lose are all the same kind of clip: fan edits that jump between scenes. The right film was actually the top candidate on every one of them, but the frames point at many different places in the film, because the clip really does jump around, so my agreement check turns them away. I measured what happens if I loosen the check to let them back in: two of the wrong answers come back with them, because on pictures alone they look the same. You told me a no answer is better than a wrong one, so I kept it strict. But that's a choice about what Trace prefers, and I'd rather you made it than me.

One more thing from the tables. There's a clip the pod answered as The Misfits, and a user voted it wrong. When I run that clip, the pod's own rules and mine both say F1, the Formula One film, at distance eight. Someone should watch that clip, because the live answer may not have come from where we think it did.

5. The edits test

Section 5. The edit table screenshot, then the before-and-after one.

For the edits test I took eight clips the engine identifies correctly, and made ten edited copies of each. The edits are the ones reposters actually do: a colour filter, black and white, re-compressing, speeding up, slowing down, zooming in, tilting it a couple of degrees, caption bars, shrinking it inside a black frame, and flipping it left-to-right.

First result, which surprised me: not one edit, on any version of the engine, produced a wrong film. When an edit beats the engine, the engine goes quiet. It doesn't guess.

Most edits barely hurt. Filters, black and white, re-compression, speed changes, nearly everything still gets identified. Two edits made it completely blind, though. The flip: zero out of eight identified. And the black frame: one out of eight.

Both turned out to be missing machinery rather than bad thresholds. Nothing anywhere in the engine ever looks at the flipped picture, so a mirrored repost can't match anything. I made the probe also fingerprint the mirror image of each frame it samples, and zero out of eight became eight out of eight, matching at distance two, basically perfect. For the black frame, the pod has a trick called the active region crop that finds the picture inside the borders, and it's one of the pieces that never made it into my laptop copy. I built a simple version: find the black bars, cut them off, fingerprint what's inside. One out of eight became six out of eight. The two still missed have bars that aren't quite black, so my simple version doesn't see them as bars.

Then I checked I hadn't broken anything. I re-ran all 125 real clips with the new views switched on, and not a single answer changed, in either direction.

The cost is speed: with the extra views the probe takes about four times longer per clip. On my laptop that doesn't matter. If this ever goes towards the pod, the way to do it is to try the normal views first and only do the mirror and border checks when the normal ones found nothing. Then an ordinary clip costs the same as it does today.

6. Housekeeping

Section 6. Quick.

Three quick things. All the code and results are in a private GitHub repo now, called pdq-testing. I can add you to it today.

Pod access works from my machine. I tested it read-only: looked at the hostname and the ready status, touched nothing.

And one thing I should flag. While I was setting that up, the access document with the key in it got pasted into a chat. If you'd rather rotate the key, this is a clean moment to do it. Your call.

7. The asks

Close the document. Count on fingers.

Four things.

One, the call on the four fan-edit clips: keep it strict, or take them back along with two wrong answers. I recommend keeping it strict.

Two, someone watches the Misfits clip.

Three, a direction on the probe. If you want it taken towards the pod, as the replacement for the pre-scan and the one-frame probe, my next step would be the normal-views-first version, measured at the pod's index size.

And four, still open from last week: keeping the clip_query_hashes rows for voted clips. The clips are still rotting.

If he asks...

...why not loosen the probe to get the four fan edits back?

Because I measured it. The four I'd get back and the two wrong ones I'd let back in sit in the same range of evidence. There's nowhere to put the bar that separates them. So it's a choice between the two, not a fix.

...whether the mirror views could cause new wrong answers?

That was my worry too, so that was the test I gated it on. All 125 real clips, before and after, zero answers changed. The diff is in the repo if you want it.

...how long the probe takes now?

About two minutes a clip on my laptop with every view on, against thirty seconds before.
Normal-views-first would bring ordinary clips back to today's cost.

...whether the speed edits broke the film-position agreement?

I expected them to, and they didn't. The check has a three-second tolerance and that absorbs a 25 percent speed change. I had a fix designed for it and threw it away, because the measurement said it wasn't needed.