



MEETING SCRIPT

Script for presenting the 4 October update

For me only. The document on the table is PDQ_WEEKLY_UPDATE_04_OCTOBER_2026.pdf.

Same rule as last time: don't read it word for word. Read a section, look up, say it the way it comes out. The numbers are the only part to get exactly right.

About eight minutes at a normal pace.

Opening

Have the document closed. Just talk.

> Last week you told me to stop bending the thresholds and change the actual tech. So > this week was three jobs. > > First, I made the testing honest: every clip I run now records which rule gave the > answer. So when something is wrong, I know exactly which part of the engine did it, > not just that it happened. > > Second, I compared the three parts of the engine that can answer, worked out what the > careful rules have that the pre-scan and the probe don't, and rebuilt the probe out of > the missing pieces. > > Third, I built an edits test. Ten things a reposter actually does to a clip, applied > one at a time to clips the engine gets right, so any failure is the edit's fault and > nothing else. That found two places where the engine is completely blind, and I fixed > both. > > The headline: the pod's rules give seven wrong answers on my full set. The rebuilt > probe with the rules behind it gives one. And of the fifty-two correct answers, I keep > forty-eight. I'll go through it.

1. Which rule gave the wrong answers

Open at section 1. Point at the screenshot.

> This is the printout, not a summary. All 125 clips through the pod's exact rules, and > every answer tagged with the rule that produced it. > > Seven wrong answers reproduce. Four of them come from the pre-scan, and look at the > distances: six, six, twelve, eighteen. Those pictures genuinely are in the index. The > rules never even ran on those clips, which is why the last two weeks of rule changes > could never have caught them. The other three are P0, P9E and P8, which are exactly > the three my earlier changes went after. > > So the thing I'd been tuning could only ever reach three of the seven. That's the > clearest argument I have for what you said last week.

2. What the system has that the deciders don't

Section 2. The small table.

> The engine can answer from three places. The pre-scan, the probe, and the rules. > > The careful rules check three things: lots of frames, the frames agreeing on where in > the film they land, and the frames

coming from more than one moment of the clip. > > The pre-scan and the probe, which between them decide almost everything, check none of > that. The probe answers off one frame. And the engine's own notes say a true match > gives the same offset for every frame. That's the strongest signal in the data and > the two busiest parts of the engine ignore it.

3. The rebuilt probe

Section 3. Don't linger, the next section is the point.

> So the new probe is just: the pre-scan's whole-clip view, plus all the checks it never > had. It samples forty frames across the clip, keeps a scorecard per film, and it only > answers on three bars, each copied from a rule the pod already trusts. Two near-perfect > frames. Or the old pre-scan bar plus agreement on film time. Or five medium frames that > agree and come from two moments of the clip. > > The old one-frame probe is retired. This replaces it.

4. The measurement

Section 4. Point at the three-line table, then the screenshot.

> Same yardstick as every week, both numbers, all 125 clips. > > Pod's rules: seven wrong, fifty-two correct. The new probe entirely on its own: one > wrong, forty-one. With the rules running behind it as the fallback: still one wrong, > and forty-eight correct. > > The one that survives is Krrish 3, at distance six. The pictures are identical to > pictures in the index. Nothing picture-based can refuse that; it needs audio or SSCD. > > The four correct answers I lose are all montage clips. The right film was the top > candidate every single time, it just scatters across film time because the clip cuts > between scenes. I can get them back, but the exact same loosening lets The Northman > and Civil War back in, because on picture evidence they look the same. You told me a > no answer beats a wrong one, so I left them out. But that's a policy call and I'd > rather you made it than me. > > One more thing from the tables. The clip the pod answered as The Misfits comes back as > F1: The Movie, at distance eight, on both the pod's rules and mine. Worth watching > that clip. Something doesn't add up about where the live answer came from.

5. The edits test

Section 5. The edit table screenshot, then the before-and-after one.

> Ten edits, each applied to eight clips the engine identifies correctly unedited. > > First thing: no edit ever produced a wrong film. Anywhere. Edits don't fool the > engine, they blind it. It goes quiet instead of guessing. > > Filters, black and white, recompression, speed changes: barely a scratch. Rotation and > zoom hurt. And two edits blinded it completely. Mirroring: zero out of eight. And > letterboxing: one out of eight. > > Both turned out to be missing tech. Nothing in the engine ever hashes a flipped frame, > so I made the probe hash the mirror of everything it samples. Zero out of eight became > eight out of eight, matching at distance two. And the pod has an active-region crop > that finds the picture inside black bars, which is one of the pieces the laptop port > left out. I built a simple version. One out of eight became six out of eight. > > Then the safety gate: I re-ran all 125 real clips with the new views on. Not one > answer changed. The extra views only add, they never subtract. > > The honest cost is the probe is about four

times slower with the extra views. On my > laptop, irrelevant. If this goes near the pod, you'd run plain views first and only > pay for the second look when the plain views find nothing.

6. Housekeeping

Section 6. Quick.

> Three small things. Everything is in a private GitHub repo now, pdq-testing, code and > results, and I can add you to it today. > > Pod access works from my machine, verified read-only, nothing touched. > > And while setting that up, the access doc with the key in it ended up pasted into a > chat. So if you'd rather rotate the key, now's a clean moment. Your call.

7. The asks

Close the document. Count on fingers.

> Four things. > > One, the policy call on the four montage clips. I recommend keeping the strict trade. > > Two, someone watches the Misfits clip. > > Three, a direction on the probe. If you want it taken towards the pod as the pre-scan > and P0 replacement, I'd start with the plain-views-first version and measure it at the > pod's index size. > > And four, still open from last week: keeping the clip_query_hashes rows for voted > clips. The clips are still rotting.

If he asks...

...why not loosen the probe to get the four montages back? > Because I tried the numbers. The four I'd recover and the two wrong ones I'd readmit > sit in the same evidence range. There's no bar between them. It's a trade, not a fix.

...whether the mirror views could cause new wrong answers? > That was my worry too, so that was the gate. All 125 real clips, before and after, > zero answers changed. It's in the repo if you want the diff.

...how long the probe takes now? > About two minutes a clip on my laptop with every view on, against thirty seconds > before. Plain-views-first brings ordinary clips back to today's cost.

...whether the speed edits broke the film-time agreement? > I expected them to and they didn't. The three-second tolerance absorbs a 25 percent > speed change. I had a fix designed and threw it away because the measurement said it > wasn't needed.